

How to learn the talk with this booklet

The booklet has three parts, each with one page per slide. Work through them in order, slide by slide, always **aloud**, standing if you can.

1. **Learn the skeleton, not the words.** Per slide, learn verbatim only the *opening sentence*, the *last sentence* (it is the bridge to the next slide) and the *numbers*. Everything in between is a list of points you say in your own words. A talk recited word for word sounds recited and collapses on the first slip; anchor sentences give you a safe place to land.
2. **Part A – full text.** One sentence per line; every line break is a pause. Read the slide twice, then cover the text and say it while looking at the thumbnail.
3. **Part B – first letters.** Say the slide from the skeleton. Peek only when stuck. Move on once you manage two clean runs.
4. **Part C – anchors only.** Say the whole slide from the opening and closing sentence alone. This is the level you need on stage.
5. **Time it.** Each page carries its time window. The pauses are part of the timing. Target for the whole talk: **9:30**.
6. **Space it out.** Day 1: slides 1–5 to Part C. Day 2: slides 6–11, then re-test 1–5. Day 3: two full runs with the PDF full-screen and a clicker, one of them recorded on your phone – listen back. Day 4: open the booklet on a random page and start from there (this trains recovery from a blank), plus the questions in QA.md. Last evening: one relaxed run. **Freeze the script two days before the talk;** every edit resets what you have learnt.
7. **Anki.** `anki.csv` holds one card per slide, one card per bridge sentence and one per number – for the minutes between sessions.

Legend. **[CLICK]** – advance the slide. *grey italics* – **[CUT IF LATE]**: can be dropped without breaking the argument. Numbers table on the last page.



Good morning.

My name is Martin Maxa, from the Department of Measurement here at this faculty.

I worked on this together with Jakub Svatoš, from the same department.

Today I will talk about acoustic detection, classification and localization of impulsive sounds, with gunshots as the target class.

The main focus is gunshot classification: we compare different cepstral features and two classifiers.

I will also briefly show the localization.

Motivation and research questions

- Acoustic detection of gunfire complements camera-based surveillance: the acoustic channel is omnidirectional and does not require line of sight.
- Three required functions, in outdoor environments that are noisy and reverberant:



detection
of the impulsive event



classification
ideally down to calibre



localization
of the source, in 2D

This study

1. Which cepstral representation separates gunshots from acoustically similar false alarms most reliably?
2. Neural network or support vector machine on those features?
3. How does the analysis frame length interact with environmental reflections?

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

1/11

This work is part of the development of an acoustic sensor system for the protection of soft targets.

Acoustic detection can complement camera-based surveillance.

Sound is omnidirectional, and it does not need a direct line of sight.

[CUT IF LATE] *The complete system should detect an impulsive event, classify it, ideally down to the calibre, and localize its source.*

It should also work outdoors, where there is noise and many reflections. [CUT]


[CLICK] _____

We focus on three questions.

First, which cepstral representation works best for separating gunshots from similar impulsive sounds?

Second, which classifier performs better: a neural network or a support vector machine?

And third, how does the frame length affect the results?



Let us first look at the signal of a gunshot.

A gunshot has one or two components, depending on the projectile speed.

A subsonic shot has only the muzzle blast; a supersonic projectile also produces a shock wave.

This recording is a supersonic shot, so we see both.

The first component, at about two and a half milliseconds, is the shock wave.

[CUT IF LATE] *It is produced along the projectile trajectory and has a characteristic N-shaped pressure profile.* [CUT]

[CLICK] _____

The second one, about forty-five milliseconds later, is the muzzle blast.

It comes from the gases at the barrel and is present in every gunshot.

Between and after these two components, we mainly see reflections from the surroundings.

They matter for the choice of frame length, and I will return to them in the results.

Single-domain analysis is insufficient

- **Spectral analysis alone** often captures environmental noise rather than the gunshot itself.
- **Time-domain patterns** get distorted by nonlinear dispersion and obstacles.

Consequence
Effective feature extraction must integrate *both* time and frequency domain processing.


Acoustic Detection of Impulsive Sounds · IMEKO TC4 20263/11

This also shows why one signal domain alone is not enough.

The spectrum often describes the environment more than the gunshot itself.

The time-domain shape changes during propagation and because of obstacles.

So we need features that combine both domains, and cepstral features do exactly that.



Now let us look at the proposed system.

It consists of several independent sensor units connected to a remote server.


Each sensor continuously monitors its surroundings.

When it detects an impulsive event, it sends the recorded signal to the server.

The server then performs two main tasks: classification of the event and localization of the sound source.

Several units give wider coverage and allow us to estimate the source position.

The system was tested with several firearm calibres and ammunition types.



First, I will briefly explain the localization.

Each sensor unit has two MEMS microphones, about twenty centimetres apart, sampled at 44.1 kilohertz.

We measure the time difference of arrival between the two microphones.

From this delay, we can calculate the direction to the source using the equation shown here.


[CLICK] _____

One sensor unit therefore gives us one direction.

With two sensor units, we get two directions, and their intersection gives us the estimated source position.

We tested this on seventy events at nine reference angles.

Using cross-correlation with parabolic interpolation, we achieved a mean absolute error of zero point seven seven degrees.




Now I will move to the main part of the work: classification.

We compared four different cepstral feature extractors.

All of them use the same chain: framing, FFT, a filter bank, and then a logarithm and a discrete cosine transform.

The main difference is the filter bank.

You can see all four filter banks here, drawn on the same frequency axis.



MFCC uses the mel scale, with dense filters at low frequencies.

It comes from speech processing, but it also works for gunshots.

LFCC uses linearly spaced filters, while IMFCC uses an inverted mel scale, with dense filters at high frequencies.

GTCC uses a gammatone filter bank based on a model of the human auditory system.

It is also reported to be more robust to noise than MFCC.

Classification results – 9 mm class							
Frame	Extractor	ACC (%)		RLC (%)		MCC (–)	
		SVM	NN	SVM	NN	SVM	NN
15 ms	MFCC	83.90	85.59	80.65	67.74	0.62	0.62
	LFCC	83.90	83.90	80.65	64.52	0.62	0.57
	IMFCC	85.59	90.68	83.87	87.10	0.66	0.77
	GTCC	93.22	97.46	87.10	93.55	0.82	0.93
30 ms	MFCC	84.75	87.29	67.74	87.10	0.60	0.70
	LFCC	87.29	83.90	70.97	70.97	0.66	0.59
	IMFCC	85.59	92.37	67.74	80.65	0.62	0.80
	GTCC	97.46	98.31	90.32	93.55	0.93	0.96
50 ms	MFCC	84.75	93.22	74.19	90.32	0.61	0.83
	LFCC	88.98	86.44	77.42	77.42	0.71	0.66
	IMFCC	89.83	90.68	74.19	80.65	0.73	0.76
	GTCC	97.46	99.15	90.32	96.77	0.93	0.98

Classification results for the 9 mm class. Test partition: 118 events, 31 of them 9 mm.


Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

8/11

This table shows the complete classification results for the nine-millimetre class.

The GTCC rows are highlighted.

I will not go through the numbers here – let us look at the same data in a chart.



This chart shows the same results, but only the Matthews correlation coefficient.

MCC considers the whole confusion matrix, so it is less affected by class imbalance than accuracy.

The red line is GTCC.

It is on top in both panels and for every frame length.

The neural network gives better results than the SVM, and the difference grows with longer frames.

The best configuration is GTCC with the neural network and a fifty-millisecond frame.

It reaches an accuracy of ninety-nine point one five percent, a recall of ninety-six point seven seven percent, and an MCC of zero point nine eight.

That is one hundred and seventeen correct decisions out of one hundred and eighteen test events.

Discussion

- MFCC, LFCC and IMFCC yield similar results.
- **GTCC consistently performs better** across every frame length and both classifiers.
- **NN outperforms SVM** – better at modelling complex patterns, especially with longer frames.
- **Longer frames improve accuracy**, but start to capture environmental reflections, which can pull focus away from the gunshot itself.

Best configuration in the table

GTCC + NN at 50 ms – 99.15 % ACC, 96.77 % RLC, MCC 0.98.

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

10/11

From these results, we can make three main observations.

First, MFCC, LFCC and IMFCC give similar results, so the filter spacing makes little difference.

Second, GTCC performs clearly better.

[CUT IF LATE] *One possible reason is that the gammatone filter bank is based on an auditory model and is known to be more robust to noise.*

However, we did not study this mechanism directly. [CUT]

Third, the neural network performs slightly better than the SVM, especially with longer frames.

This is a general trend, not true for every single result.

[CLICK] _____

Finally, longer frames generally improve the classification accuracy.

But with a fifty-millisecond frame, we already start to include more reflections from the environment, not only the gunshot itself.

So even longer frames would not necessarily help.

Conclusion

- A distributed network of stand-alone acoustic sensors with a remote server detects, classifies and localizes impulsive sounds in 2D.
- Cepstral coefficients are an effective feature extractor for gunshot classification with both NN and SVM.
- GTCC beats Mel-frequency based coefficients; NN achieves slightly better classification success, especially with longer frames.
- Validated with various firearm calibers and ammunition.

Thank you – questions?

svatoja1@fel.cvut.cz maxamart@fel.cvut.cz

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

11/11

To conclude.

We presented a distributed acoustic sensor system for the detection, classification and localization of impulsive sounds.

For gunshot classification, cepstral features work well with both tested classifiers.

Among the tested features, GTCC gives the best results.

The neural network also performs slightly better than the SVM, especially with longer frames.

The system was evaluated on real firearms with several calibres and ammunition types.

Thank you.



G m.

M n i M M, f t D o M h a t f.

I w o t t w J S, f t s d.

T I w t a a d, c a l o i s, w g a t t c.

T m f i g c: w c d c f a t c.

I w a b s t l.

Motivation and research questions

- Acoustic detection of gunfire complements camera-based surveillance: the acoustic channel is omnidirectional and does not require line of sight.
- Three required functions, in outdoor environments that are noisy and reverberant:



detection
of the impulsive event



classification
ideally down to calibre



localization
of the source, in 2D

This study

- Which cepstral representation separates gunshots from acoustically similar false alarms most reliably?
- Neural network or support vector machine on those features?
- How does the analysis frame length interact with environmental reflections?

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

1/11

T w i p o t d o a a s s f t p o s t .

A d c c c - b s .

S i o , a i d n n a d l o s .

T c s s d a i e , c i , i d t t c , a l i s .

I s a w o , w t i n a m r .


[CLICK]

W f o t q .

F , w c r w b f s g f s i s ?

S , w c p b : a n n o a s v m ?

A t , h d t f l a t r ?



L u f l a t s o a g.

A g h o o t c, d o t p s.

A s s h o t m b; a s p a p a s w.

T r i a s s, s w s b.

T f c, a a t a a h m, i t s w.

I i p a t p t a h a c N-s p p.

[CLICK]

T s o, a f-f m l, i t m b.

I c f t g a t b a i p i e g.

B a a t t c, w m s r f t s.

T m f t c o f l, a I w r t t i t r.

Single-domain analysis is insufficient

- **Spectral analysis alone** often captures environmental noise rather than the gunshot itself.
- **Time-domain patterns** get distorted by nonlinear dispersion and obstacles.

Consequence

Effective feature extraction must integrate *both* time and frequency domain processing.

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026


3/11

T a s w o s d a i n e.

T s o d t e m t t g i.

T t-d s c d p a b o o.

S w n f t c b d, a c f d e t.



N l u l a t p s.

I c o s i s u c t a r s.


E s c m i s.

W i d a i e, i s t r s t t s.

T s t p t m t: c o t e a l o t s s.

S u g w c a a u t e t s p.

T s w t w s f c a a t.



F, I w b e t l.

E s u h t MEMS m, a t c a, s a 44.1 k.

W m t t d o a b t t m.

F t d, w c c t d t t s u t e s h.


[CLICK] _____

O s u t g u o d.

W t s u, w g t d, a t i g u t e s p.

W t t o s e a n r a.

U c-c w p i, w a a m a e o z p s s d.



N I w m t t m p o t w: c.

W c f d c f e.

A o t u t s c: f, FFT, a f b, a t a l a a d c t.

T m d i t f b.

Y c s a f f b h, d o t s f a.

Feature extraction – the four filter banks



MFCC – mel scale
dense at low frequencies;
speech-born, validated on gunshots at military ranges



LFCC – linear
equally spaced filters, uniform resolution across the band



IMFCC – inverted mel
mirror of the mel scale: dense at high frequencies



GTCC – gammatone
auditory model, ERB spacing; **reported more noise-robust than MFCC**

Experimental grid: 4 extractors × 3 frame lengths (15 / 30 / 50 ms) × 2 classifiers (NN, SVM).

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026 7/11

MFCC u t m s, w d f a l f.

I c f s p, b i a w f g.

LFCC u l s f, w IMFCC u a i m s, w d f a h f.

GTCC u a g f b b o a m o t h a s.

I i a r t b m r t n t MFCC.

Classification results – 9 mm class							
Frame	Extractor	ACC (%)		RLC (%)		MCC (–)	
		SVM	NN	SVM	NN	SVM	NN
15 ms	MFCC	83.90	85.59	80.65	67.74	0.62	0.62
	LFCC	83.90	83.90	80.65	64.52	0.62	0.57
	IMFCC	85.59	90.68	83.87	87.10	0.66	0.77
	GTCC	93.22	97.46	87.10	93.55	0.82	0.93
30 ms	MFCC	84.75	87.29	67.74	87.10	0.60	0.70
	LFCC	87.29	83.90	70.97	70.97	0.66	0.59
	IMFCC	85.59	92.37	67.74	80.65	0.62	0.80
	GTCC	97.46	98.31	90.32	93.55	0.93	0.96
50 ms	MFCC	84.75	93.22	74.19	90.32	0.61	0.83
	LFCC	88.98	86.44	77.42	77.42	0.71	0.66
	IMFCC	89.83	90.68	74.19	80.65	0.73	0.76
	GTCC	97.46	99.15	90.32	96.77	0.93	0.98


Classification results for the 9 mm class. Test partition: 118 events, 31 of them 9 mm.

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026 8/11

T t s t c c r f t n-m c.

T GTCC r a h.

I w n g t t n h - l u l a t s d i a c.



T c s t s r, b o t M c c.

MCC c t w c m, s i i l a b c i t a.

T r l i GTCC.

I i o t i b p a f e f l.

T n n g b r t t SVM, a t d g w l f.

T b c i GTCC w t n n a a f-m f.

I r a a o n-n p o f p, a r o n-s p s s p, a a MCC o z p n e.

T i o h a s c d o o o h a e t e.

Discussion

- MFCC, LFCC and IMFCC yield similar results.
- **GTCC consistently performs better** across every frame length and both classifiers.
- **NN outperforms SVM** – better at modelling complex patterns, especially with longer frames.
- **Longer frames improve accuracy**, but start to capture environmental reflections, which can pull focus away from the gunshot itself.

Best configuration in the table

GTCC + NN at 50 ms – 99.15 % ACC, 96.77 % RLC, MCC 0.98.

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

10/11

F t r, w c m t m o.

F, MFCC, LFCC a IMFCC g s r, s t f s m l d.

S, GTCC p c b.

O p r i t t g f b i b o a a m a i k t b m r t n.

H, w d n s t m d.

T, t n n p s b t t SVM, e w l f.

T i a g t, n t f e s r.

[CLICK] _____

F, l f g i t c a.

B w a f-m f, w a s t i m r f t e, n o t g i.

S e l f w n n h.

Conclusion

- A distributed network of stand-alone acoustic sensors with a remote server detects, classifies and localizes impulsive sounds in 2D.
- Cepstral coefficients are an effective feature extractor for gunshot classification with both NN and SVM.
- GTCC beats Mel-frequency based coefficients; NN achieves slightly better classification success, especially with longer frames.
- Validated with various firearm calibers and ammunition.

Thank you – questions?

svatoja1@fel.cvut.cz maxamart@fel.cvut.cz

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

11/11

T c.

W p a d a s s f t d, c a l o i s.

F g c, c f w w w b t c.

A t t f, GTCC g t b r.

T n n a p s b t t SVM, e w l f.

T s w e o r f w s c a a t.

T y.



Good morning.

...

I will also briefly show the localization.

Motivation and research questions

- Acoustic detection of gunfire complements camera-based surveillance: the acoustic channel is omnidirectional and does not require line of sight.
- Three required functions, in outdoor environments that are noisy and reverberant:



detection
of the impulsive event



classification
ideally down to calibre



localization
of the source, in 2D

This study

1. Which cepstral representation separates gunshots from acoustically similar false alarms most reliably?
2. Neural network or support vector machine on those features?
3. How does the analysis frame length interact with environmental reflections?


Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

1/11

This work is part of the development of an acoustic sensor system for the protection of soft targets.

• • •

And third, how does the frame length affect the results?



Let us first look at the signal of a gunshot.

...

They matter for the choice of frame length, and I will return to them in the results.

Single-domain analysis is insufficient

- **Spectral analysis alone** often captures environmental noise rather than the gunshot itself.
- **Time-domain patterns** get distorted by nonlinear dispersion and obstacles.

Consequence


Effective feature extraction must integrate *both* time and frequency domain processing.

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026 3/11

This also shows why one signal domain alone is not enough.

• • •


So we need features that combine both domains, and cepstral features do exactly that.



Now let us look at the proposed system.

...


The system was tested with several firearm calibres and ammunition types.



First, I will briefly explain the localization.

...


Using cross-correlation with parabolic interpolation, we achieved a mean absolute error of zero point seven seven degrees.



Now I will move to the main part of the work: classification.

...

You can see all four filter banks here, drawn on the same frequency axis.



MFCC uses the mel scale, with dense filters at low frequencies.

...

It is also reported to be more robust to noise than MFCC.

Classification results – 9 mm class							
Frame	Extractor	ACC (%)		RLC (%)		MCC (–)	
		SVM	NN	SVM	NN	SVM	NN
15 ms	MFCC	83.90	85.59	80.65	67.74	0.62	0.62
	LFCC	83.90	83.90	80.65	64.52	0.62	0.57
	IMFCC	85.59	90.68	83.87	87.10	0.66	0.77
	GTCC	93.22	97.46	87.10	93.55	0.82	0.93
30 ms	MFCC	84.75	87.29	67.74	87.10	0.60	0.70
	LFCC	87.29	83.90	70.97	70.97	0.66	0.59
	IMFCC	85.59	92.37	67.74	80.65	0.62	0.80
	GTCC	97.46	98.31	90.32	93.55	0.93	0.96
50 ms	MFCC	84.75	93.22	74.19	90.32	0.61	0.83
	LFCC	88.98	86.44	77.42	77.42	0.71	0.66
	IMFCC	89.83	90.68	74.19	80.65	0.73	0.76
	GTCC	97.46	99.15	90.32	96.77	0.93	0.98


Classification results for the 9 mm class. Test partition: 118 events, 31 of them 9 mm.

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026 8/11

This table shows the complete classification results for the nine-millimetre class.

• • •

I will not go through the numbers here – let us look at the same data in a chart.



This chart shows the same results, but only the Matthews correlation coefficient.

...

That is one hundred and seventeen correct decisions out of one hundred and eighteen test events.

Discussion

- MFCC, LFCC and IMFCC yield similar results.
- **GTCC consistently performs better** across every frame length and both classifiers.
- **NN outperforms SVM** – better at modelling complex patterns, especially with longer frames.
- **Longer frames improve accuracy**, but start to capture environmental reflections, which can pull focus away from the gunshot itself.

Best configuration in the table

GTCC + NN at 50 ms – 99.15 % ACC, 96.77 % RLC, MCC 0.98.

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

10/11

From these results, we can make three main observations.

• • •

So even longer frames would not necessarily help.

Conclusion

- A distributed network of stand-alone acoustic sensors with a remote server detects, classifies and localizes impulsive sounds in 2D.
- Cepstral coefficients are an effective feature extractor for gunshot classification with both NN and SVM.
- GTCC beats Mel-frequency based coefficients; NN achieves slightly better classification success, especially with longer frames.
- Validated with various firearm calibers and ammunition.

Thank you – questions?

svatoja1@fel.cvut.cz maxamart@fel.cvut.cz

Acoustic Detection of Impulsive Sounds · IMEKO TC4 2026

11/11

To conclude.

• • •

Thank you.

Every number in the talk

Best configuration (GTCC + NN, 50 ms) – accuracy?	99.15 %
Best configuration (GTCC + NN, 50 ms) – recall?	96.77 %
Best configuration (GTCC + NN, 50 ms) – MCC?	0.98
Test partition – how many events, how many 9 mm?	118 events, 31 of them 9 mm
99.15 % accuracy = how many correct decisions?	117 of 118
96.77 % recall = how many 9 mm shots recalled?	30 of 31
Corpus – training / test samples (five classes)?	371 / 118
Features per extractor?	26 coefficients per frame
Frame lengths tested?	15 / 30 / 50 ms
Bearing accuracy (cross-correlation + parabolic interpolation)?	MAE 0.77°
Bearing evaluation – how many events, how many reference angles?	70 events, 9 angles (−45° to +60°)
Microphone spacing and sampling rate?	d = 0.186 m (≈ 20 cm), f_s = 44.1 kHz
Maximum TDoA for this geometry?	≈ 24 samples
Shock wave / muzzle blast timing in the recording shown?	shock wave ≈ 2.5 ms, muzzle blast ≈ 45 ms later, lasting 1–3 ms
Talk length – target?	10 min slot, script written to 9:30
